

# HW4

Siddharth Agarwal

19/04/2021

load packages

```
if(!require('pacman')) install.packages('pacman')

## Loading required package: pacman

pacman::p_load(ISLR, MASS, rpart, rpart.plot, caret, leaps, randomForest,
               gbm, tree, ggplot2, dplyr, tinytex)

#tinytex::install_tinytex()
```

**Q1. Remove the observations with unknown salary information.**

How many observations were removed in this process?

```
data('Hitters')

summary(Hitters$Salary)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	67.5	190.0	425.0	535.9	750.0	2460.0	59

```
Hitters.modified <- subset(Hitters, !is.na(Hitters$Salary))


#summary(Hitters.modified)
```

**Ans 1:** 59 observations having unknown salary were removed

**Q2. Transform the salaries using a (natural) log transformation.**


Is there any justification for this transformation? Explain your answer.

```
boxplot(Hitters.modified$Salary, col = 'green', main='Salary Boxplot')
```



```
Hitters.modified <- transform(Hitters.modified,  
                              log.salary = log(Hitters.modified$Salary))  
  
boxplot(Hitters.modified$log.salary, col = 'green', main='Salary Boxplot')
```

## Salary Boxplot



```
head(Hitters.modified)
```


```
##           AtBat Hits HmRun Runs RBI Walks Years CAtBat CHits CHmRun
## -Alan Ashby    315   81    7  24  38   39   14   3449   835    69
## -Alvin Davis   479  130   18  66  72   76    3   1624   457    63
## -Andre Dawson  496  141   20  65  78   37   11   5628  1575   225
## -Andres Galarraga 321   87   10  39  42   30    2    396   101    12
## -Alfredo Griffin 594  169    4  74  51   35   11   4408  1133    19
## -Al Newman    185   37    1  23   8   21    2    214    42     1
##           CRuns CRBI CWalks League Division PutOuts Assists Errors
## -Alan Ashby    321  414   375     N         W      632    43     10
## -Alvin Davis   224  266   263     A         W      880    82     14
## -Andre Dawson  828  838   354     N         E     200    11      3
## -Andres Galarraga 48   46    33     N         E     805    40      4
## -Alfredo Griffin 501  336   194     A         W     282   421     25
## -Al Newman     30    9    24     N         E      76   127      7
##           Salary NewLeague log.salary
## -Alan Ashby    475.0         N   6.163315
## -Alvin Davis   480.0         A   6.173786
## -Andre Dawson  500.0         N   6.214608
## -Andres Galarraga 91.5         N   4.516339
## -Alfredo Griffin 750.0         A   6.620073
## -Al Newman     70.0         A   4.248495
```

Ans 2: The data for Salary column is positively skewed as all the outlier

values have values greater than 1500. Hence it would be better to perform a log transformation here to get the data to be more normal in its distribution

Q3. Create a scatterplot with Hits on the y-axis and Years on the x-axis using all the observations. Color code the observations using the log Salary variable. What patterns do you notice on this chart, if any?

```
ggplot(Hitters.modified, aes(y=Hits, x=Years, color= log.salary)) +  
  geom_point(alpha=0.8)
```



Ans 3: The graph shows that Salary tends to increase as the number of Years or Hits increase

We can see that there are few very low salaries ( $\text{log.salary} \sim 5$  or less) for Years  $> 5$ .

For Hits we do not observe any particular trend for salary values especially for Years of experience  $< 5$  as low salaries are distributed throughout the vertical axis.

There are few very high salaries (outliers) in the graph. One particular surprising observation is a very high salary for (Years = 2 and Hits = 0)

Q4. Run a linear regression model of Log Salary on all the predictors using the entire dataset. Use `regsubsets()` function to perform best subset selection from the regression model.

Identify the best model using BIC.

Which predictor variables are included in this (best) model?

```
hitter.lm <- lm(log.salary ~. -Salary, Hitters.modified)
summary(hitter.lm)
```

```
##
## Call:
## lm(formula = log.salary ~ . - Salary, data = Hitters.modified)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-2.22870	-0.45350	0.09424	0.40474	2.77223

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	4.618e+00	1.765e-01	26.171	< 2e-16 ***
AtBat	-2.984e-03	1.232e-03	-2.421	0.01620 *
Hits	1.308e-02	4.622e-03	2.831	0.00503 **
HmRun	1.179e-02	1.205e-02	0.978	0.32889
Runs	-1.419e-03	5.794e-03	-0.245	0.80670
RBI	-1.675e-03	5.056e-03	-0.331	0.74063
Walks	1.096e-02	3.554e-03	3.082	0.00229 **
Years	5.696e-02	2.413e-02	2.361	0.01902 *
CAtBat	1.283e-04	2.629e-04	0.488	0.62596
CHits	-4.414e-04	1.311e-03	-0.337	0.73670
CHmRun	-7.809e-05	3.144e-03	-0.025	0.98020
CRuns	1.513e-03	1.459e-03	1.037	0.30072
CRBI	1.312e-04	1.346e-03	0.097	0.92246
CWalks	-1.466e-03	6.377e-04	-2.298	0.02239 *
LeagueN	2.825e-01	1.541e-01	1.833	0.06797 .
DivisionW	-1.656e-01	7.847e-02	-2.111	0.03580 *
PutOuts	3.389e-04	1.505e-04	2.251	0.02526 *
Assists	6.214e-04	4.300e-04	1.445	0.14970
Errors	-1.197e-02	8.537e-03	-1.402	0.16225

```
## NewLeagueN -1.742e-01 1.536e-01 -1.134 0.25788
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6135 on 243 degrees of freedom
## Multiple R-squared:  0.5586, Adjusted R-squared:  0.524
## F-statistic: 16.18 on 19 and 243 DF,  p-value: < 2.2e-16
```

```
#BIC
```

```
set.seed(42)
```

```
hitter.BIC <- regsubsets(log.salary~.-Salary, data = Hitters.modified,
                        nbest = 1, nvmax = 19, method = 'seq')
```

```
sum <- summary(hitter.BIC)
```

```
sum$which
```

```
##      (Intercept) AtBat  Hits HmRun  Runs   RBI Walks Years CAtBat CHits CHmRun
## 1             TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 2             TRUE FALSE FALSE FALSE TRUE  FALSE FALSE FALSE FALSE TRUE  FALSE
## 3             TRUE FALSE TRUE  FALSE FALSE FALSE TRUE  FALSE FALSE TRUE  FALSE
## 4             TRUE TRUE  TRUE  FALSE FALSE TRUE  FALSE TRUE  FALSE TRUE  FALSE
## 5             TRUE TRUE  TRUE  TRUE  TRUE TRUE  FALSE FALSE FALSE FALSE FALSE
## 6             TRUE TRUE  TRUE  FALSE FALSE TRUE  TRUE  TRUE  FALSE TRUE  FALSE
## 7             TRUE TRUE  TRUE  FALSE FALSE TRUE  TRUE  TRUE  FALSE TRUE  FALSE
## 8             TRUE TRUE  TRUE  FALSE FALSE TRUE  TRUE  TRUE  FALSE FALSE FALSE
## 9             TRUE TRUE  TRUE  FALSE FALSE TRUE  TRUE  TRUE  FALSE FALSE FALSE
## 10            TRUE TRUE  TRUE  FALSE FALSE TRUE  TRUE  TRUE  FALSE FALSE FALSE
## 11            TRUE TRUE  TRUE  TRUE  FALSE FALSE TRUE  TRUE  TRUE  FALSE FALSE
## 12            TRUE TRUE  TRUE  TRUE  FALSE FALSE TRUE  TRUE  TRUE  FALSE FALSE
## 13            TRUE TRUE  TRUE  TRUE  TRUE  TRUE TRUE  TRUE  TRUE  TRUE  TRUE
## 14            TRUE TRUE  TRUE  TRUE  TRUE  TRUE TRUE  TRUE  TRUE  TRUE  TRUE
## 15            TRUE TRUE  TRUE  TRUE  TRUE  TRUE TRUE  TRUE  TRUE  TRUE  TRUE
## 16            TRUE TRUE  TRUE  TRUE  FALSE TRUE  TRUE  TRUE  TRUE  TRUE  FALSE
## 17            TRUE TRUE  TRUE  TRUE  TRUE  TRUE TRUE  TRUE  TRUE  TRUE  TRUE
## 18            TRUE TRUE  TRUE  TRUE  TRUE  TRUE TRUE  TRUE  TRUE  TRUE  FALSE
## 19            TRUE TRUE  TRUE  TRUE  TRUE  TRUE TRUE  TRUE  TRUE  TRUE  TRUE
##      CRuns  CRBI CWalks LeagueN DivisionW PutOuts Assists Errors NewLeagueN
## 1      TRUE FALSE FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE
## 2     FALSE FALSE FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE
## 3     FALSE FALSE FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE
## 4     FALSE FALSE FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE
## 5     FALSE FALSE FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE
## 6     FALSE FALSE FALSE  FALSE  TRUE  FALSE  FALSE  FALSE  FALSE  FALSE
## 7     FALSE FALSE FALSE  FALSE  TRUE  TRUE  FALSE  FALSE  FALSE  FALSE
## 8      TRUE FALSE TRUE  FALSE  TRUE  TRUE  FALSE  FALSE  FALSE  FALSE
## 9      TRUE FALSE TRUE  TRUE  TRUE  TRUE  FALSE  FALSE  FALSE  FALSE
## 10     TRUE FALSE TRUE  TRUE  TRUE  TRUE  FALSE  FALSE  TRUE  TRUE
## 11     TRUE FALSE TRUE  TRUE  TRUE  TRUE  FALSE  FALSE  TRUE  TRUE
## 12     TRUE FALSE TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  FALSE
## 13     TRUE TRUE  TRUE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE
## 14     TRUE TRUE  TRUE  TRUE  FALSE  FALSE  FALSE  FALSE  FALSE  FALSE
```

```
## 15 TRUE TRUE TRUE TRUE TRUE FALSE FALSE FALSE FALSE
## 16 TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## 17 TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE FALSE
## 18 TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## 19 TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

```
sum$bic
```

```
## [1] -117.03045 -156.35434 -158.14586 -159.21816 -36.91193 -157.92069
## [7] -156.97937 -156.19540 -152.76488 -148.80615 -144.59624 -140.65413
## [13] -120.32941 -118.07255 -117.47646 -120.19950 -111.48262 -109.18595
## [19] -103.61446
```

In the best model as determined by lowest value of BIC,

the variables ‘AtBat’, ‘Hits’, ‘Walks’ and ‘CAtBat’ should be included

Q5. Now create a training data set consisting of 80 percent of the observations, and a test data set consisting of the remaining 20 percent of the observations.

```
set.seed(42)

train <- sample(1:nrow(Hitters.modified), round(nrow(Hitters.modified)*0.8))
train.data <- Hitters.modified[train,]
test.data <- Hitters.modified[-train,]
```

Q6. Generate a regression tree of log Salary using only Years and Hits variables from the training data set. Which players are likely to receive highest salaries according to this model? Write down the rule and elaborate on it.

```
#regression

tree.reg <- rpart(log.salary ~ Years+Hits, data = train.data, cp= 0.01,
                  minsplit=10, xval=10)

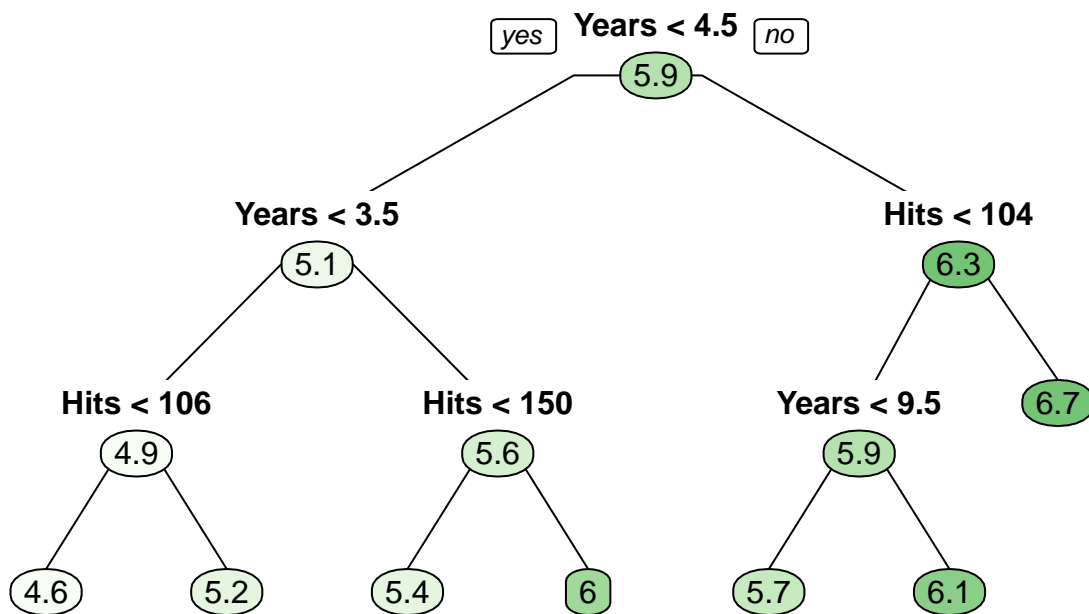
printcp(tree.reg)

##
## Regression tree:
## rpart(formula = log.salary ~ Years + Hits, data = train.data,
##       cp = 0.01, minsplit = 10, xval = 10)
##
```

```
## Variables actually used in tree construction:
## [1] Hits Years
##
## Root node error: 159.02/210 = 0.75722
##
## n= 210
##
##      CP nsplit rel error  xerror   xstd
## 1 0.454751    0  1.00000 1.01102 0.075701
## 2 0.139001    1  0.54525 0.55093 0.055952
## 3 0.050758    2  0.40625 0.41578 0.050642
## 4 0.021216    3  0.35549 0.36725 0.049628
## 5 0.012284    4  0.33427 0.37388 0.057901
## 6 0.011025    5  0.32199 0.38741 0.058225
## 7 0.010000    6  0.31097 0.38115 0.058147
```

```
# plotting tree
```

```
prp(tree.reg, type = 1, under= TRUE, roundint = FALSE, split.font = 2,
     varlen = -10, box.palette = 'Green')
```



```
# rules
```

```
rpart.rules(tree.reg, cover = TRUE)
```

```
## log.salary
```

```
cover
```



##	4.6	when Years < 4	& Hits < 106	14%
##	5.2	when Years < 4	& Hits >= 106	9%
##	5.4	when Years is 4 to 5	& Hits < 150	8%
##	5.7	when Years is 5 to 10	& Hits < 104	15%
##	6.0	when Years is 4 to 5	& Hits >= 150	3%
##	6.1	when Years >= 10	& Hits < 104	15%
##	6.7	when Years >= 5	& Hits >= 104	37%

Rule:  $\log(\text{salary}) = 6.7$  for Years  $\geq 5$  AND Hits  $\geq 104$ . The players who receive the highest salary are the one who have played for 5 years or more and made 104 hits or more. These players get receive a salary of an average 6.7 log salary which is almost equivalent to 812K in salary and it covers 37% of data

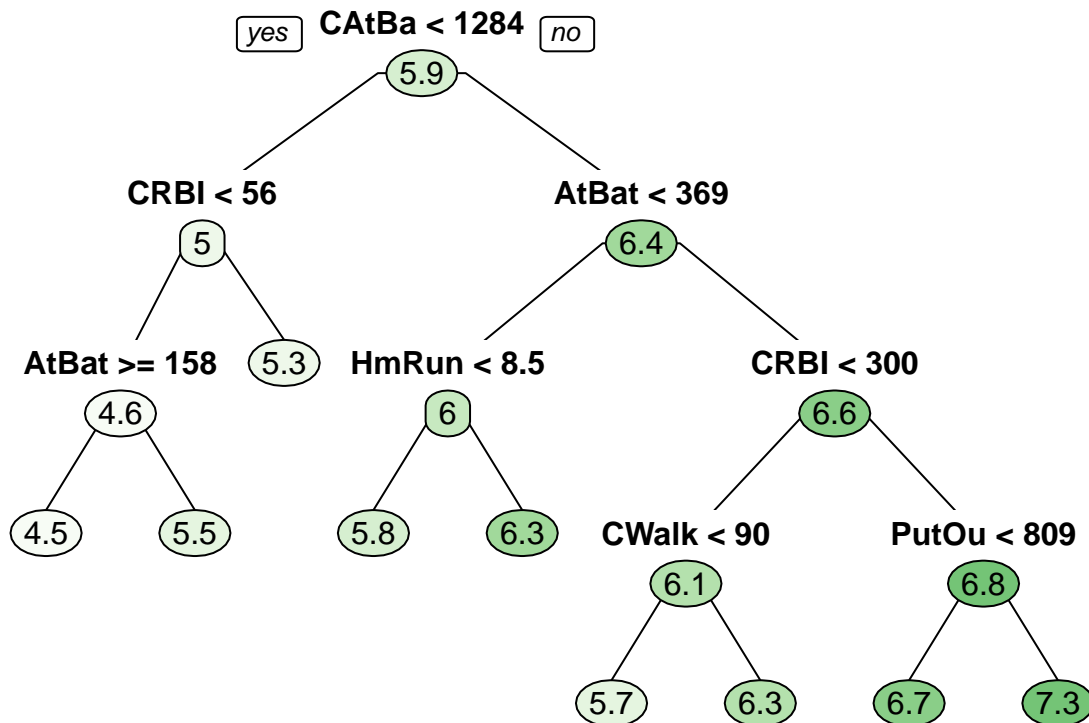
**Q7** Now create a regression tree using all the variables in the training data set. Perform boosting on the training set with 1,000 trees for a range of values of the shrinkage parameter.

Produce a plot with different shrinkage values on the x-axis and the corresponding training set MSE on the y-axis.

```
set.seed(42)

reg.all <- rpart(log.salary~.-Salary, data = train.data, cp=0.01, minsplit=10,
                xval=10)

prp(reg.all, type=1, under=TRUE, roundint=FALSE, split.font=2, varlen=-5,
     box.palette = 'Green')
```



*# Shrinkage parameter and MSE*

```
lambda <- seq(0.001, 0.1, by=0.01)
MSE <- rep(NA, length(lambda))
```

*#Boosting*

```
for(i in 1: length(lambda))
{
  boost.salary <- gbm(log.salary~.-Salary, data = train.data, n.trees = 1000,
    distribution = "gaussian", shrinkage = lambda[i])

  predictions <- predict(boost.salary, train.data, n.trees=1000,
    shrinkage=lambda[i])

  MSE[i] <- mean((predictions - train.data$log.salary)^2)
}


y <- data.frame(MSE, lambda)

y
```

```
##           MSE lambda
## 1  0.34174773  0.001
## 2  0.12305899  0.011
## 3  0.09798219  0.021
```

```
## 4 0.08399983 0.031
## 5 0.07485064 0.041
## 6 0.06724112 0.051
## 7 0.05780616 0.061
## 8 0.05141976 0.071
## 9 0.04724223 0.081
## 10 0.04556911 0.091
```

```
plot(lambda, MSE, pch=19, col='blue', type = 'b', xlab = 'Shrinkage Values',
      ylab='Mean Squared Errors',
      main = 'Shrinkage values VS MSE of training data')
```



### We can observe that as Shrinkage value increase MSE value decreases.

**Q8** Produce a plot with different shrinkage values on the x-axis and the corresponding test set MSE on the y-axis.

```
set.seed(42)

MSE2 <- rep(NA, length(lambda))

for(i in 1: length(lambda))
{
```

```

boost.test <- gbm(log.salary~.-Salary, data = train.data, n.trees = 1000,
                  distribution = "gaussian", shrinkage = lambda[i])

predictions.test <- predict(boost.test, test.data, n.trees=1000)

MSE2[i] <- mean((predictions.test - test.data$log.salary)^2)
}

y1 <- data.frame(MSE2, lambda)

y1

```

```

##           MSE2 lambda
## 1  0.4678175  0.001
## 2  0.3522949  0.011
## 3  0.3396510  0.021
## 4  0.3356590  0.031
## 5  0.3401854  0.041
## 6  0.3358513  0.051
## 7  0.3208693  0.061
## 8  0.3447181  0.071
## 9  0.3458967  0.081
## 10 0.3448906  0.091

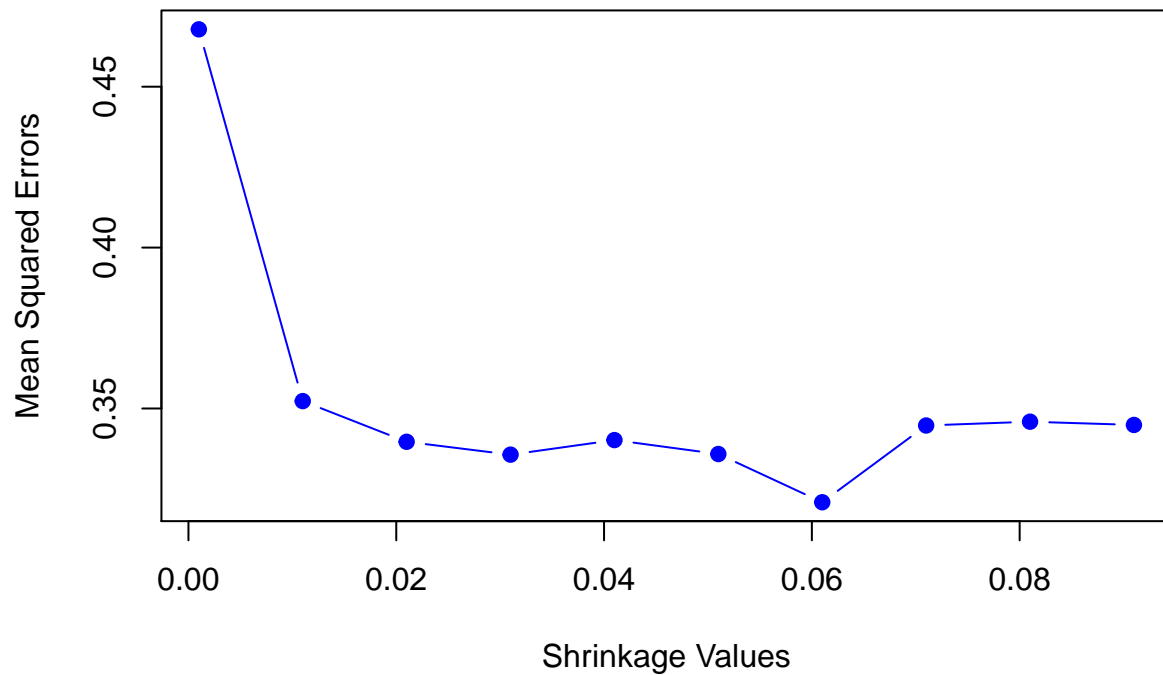
```

```

plot(lambda, MSE2, pch=19, col='blue', type = 'b', xlab = 'Shrinkage Values',
      ylab='Mean Squared Errors', main = 'Shrinkage values VS MSE of test data')

```

## Shrinkage values VS MSE of test data



### We can observe that as the lambda value increases the MSE values decrease. ### lowest value of MSE is observed at lambda value 0.061

**Q9 Which variables appear to be the most important predictors in the boosted model?**


```
set.seed(42)

# Boosting using best shrinkage value from last question

boost.train <- gbm(log.salary~. - Salary, data = train.data, n.trees = 1000,
                  shrinkage = 0.061)
```

## Distribution not specified, assuming gaussian ...

```
summary(boost.train)
```



##	var	rel.inf
## CATBat	CATBat	19.6118993
## CRBI	CRBI	8.5521956
## Years	Years	8.4482322
## PutOuts	PutOuts	7.4966800
## CRuns	CRuns	7.2955661
## CHmRun	CHmRun	6.2845306
## CWalks	CWalks	6.2605352
## CHits	CHits	5.6194432
## Hits	Hits	5.1592673
## Walks	Walks	4.9827803
## HmRun	HmRun	3.9833972
## AtBat	AtBat	3.9501757
## RBI	RBI	3.8532987
## Errors	Errors	2.9726527
## Runs	Runs	2.6922767
## Assists	Assists	1.9540525
## League	League	0.4171273
## Division	Division	0.2716072
## NewLeague	NewLeague	0.1942821

The most important predictors for the boosted model are CATBat, CRBI and CRuns

**Q10. Now apply bagging to the training set.**

**What is the test set MSE for this approach?**

```
set.seed(42)

bag <- randomForest(log.salary~. - Salary, data = train.data, mtry=19,
                    importance=TRUE)
bag.pred <- predict(bag, test.data)

mse.bag <- mean((bag.pred- test.data$log.salary)^2)

mse.bag
```

```
## [1] 0.2463472
```

**Test MSE value for bagging approach is 0.2463472**